

```
In [1]: import pandas as pd
import numpy as np
file_path= r"C:\Users\niruw\OneDrive\Desktop\Project\Placement_Data_Full_Class.csv"
df=pd.read_csv(file_path)
df.head()
```

```
Out[1]:
```

	sl_no	gender	ssc_p	ssc_b	hsc_p	hsc_b	hsc_s	degree_p	degree_t	worke
0	1	M	67.00	Others	91.00	Others	Commerce	58.00	Sci&Tech	N
1	2	M	79.33	Central	78.33	Others	Science	77.48	Sci&Tech	Y
2	3	M	65.00	Central	68.00	Central	Arts	64.00	Comm&Mgmt	N
3	4	M	56.00	Central	52.00	Central	Science	52.00	Sci&Tech	N
4	5	M	85.80	Central	73.60	Central	Commerce	73.30	Comm&Mgmt	N

```
In [2]: #data exploration
# Check basic structure
print(df.shape)
print(df.columns)
print(df.info())
print(df.describe())
print(df.head())
```

```
(215, 15)
Index(['sl_no', 'gender', 'ssc_p', 'ssc_b', 'hsc_p', 'hsc_b', 'hsc_s',
      'degree_p', 'degree_t', 'workex', 'etest_p', 'specialisation', 'mba_p',
      'status', 'salary'],
      dtype='object')
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 215 entries, 0 to 214
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	sl_no	215 non-null	int64
1	gender	215 non-null	object
2	ssc_p	215 non-null	float64
3	ssc_b	215 non-null	object
4	hsc_p	215 non-null	float64
5	hsc_b	215 non-null	object
6	hsc_s	215 non-null	object
7	degree_p	215 non-null	float64
8	degree_t	215 non-null	object
9	workex	215 non-null	object
10	etest_p	215 non-null	float64
11	specialisation	215 non-null	object
12	mba_p	215 non-null	float64
13	status	215 non-null	object
14	salary	148 non-null	float64

```
dtypes: float64(6), int64(1), object(8)
```

```
memory usage: 25.3+ KB
```

```
None
```

	sl_no	ssc_p	hsc_p	degree_p	etest_p	mba_p	\
count	215.000000	215.000000	215.000000	215.000000	215.000000	215.000000	
mean	108.000000	67.303395	66.333163	66.370186	72.100558	62.278186	
std	62.209324	10.827205	10.897509	7.358743	13.275956	5.833385	
min	1.000000	40.890000	37.000000	50.000000	50.000000	51.210000	
25%	54.500000	60.600000	60.900000	61.000000	60.000000	57.945000	
50%	108.000000	67.000000	65.000000	66.000000	71.000000	62.000000	
75%	161.500000	75.700000	73.000000	72.000000	83.500000	66.255000	
max	215.000000	89.400000	97.700000	91.000000	98.000000	77.890000	

	salary
count	148.000000
mean	288655.405405
std	93457.452420
min	200000.000000
25%	240000.000000
50%	265000.000000
75%	300000.000000
max	940000.000000

	sl_no	gender	ssc_p	ssc_b	hsc_p	hsc_b	hsc_s	degree_p	\
0	1	M	67.00	Others	91.00	Others	Commerce	58.00	
1	2	M	79.33	Central	78.33	Others	Science	77.48	
2	3	M	65.00	Central	68.00	Central	Arts	64.00	
3	4	M	56.00	Central	52.00	Central	Science	52.00	
4	5	M	85.80	Central	73.60	Central	Commerce	73.30	

	degree_t	workex	etest_p	specialisation	mba_p	status	salary
0	Sci&Tech	No	55.0	Mkt&HR	58.80	Placed	270000.0

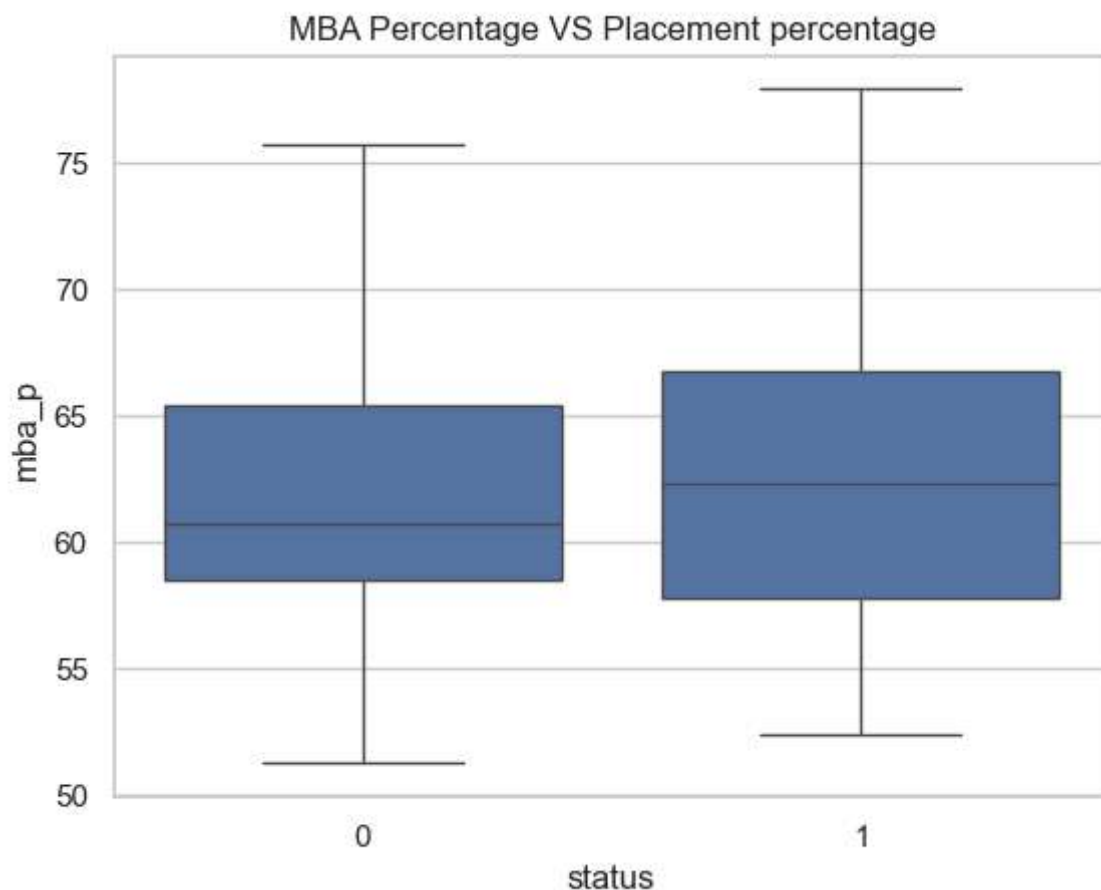
1	Sci&Tech	Yes	86.5	Mkt&Fin	66.28	Placed	200000.0
2	Comm&Mgmt	No	75.0	Mkt&Fin	57.80	Placed	250000.0
3	Sci&Tech	No	66.0	Mkt&HR	59.43	Not Placed	NaN
4	Comm&Mgmt	No	96.8	Mkt&Fin	55.50	Placed	425000.0

```
In [8]: import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style="whitegrid")
```

```
In [9]: #Data Preprocessing
df.columns = df.columns.str.strip()
df.drop(columns=['sl_no', 'salary'], inplace=True, errors='ignore')


from sklearn.preprocessing import LabelEncoder
for column in df.select_dtypes(include='object').columns:
    df[column]=LabelEncoder().fit_transform(df[column])
```

```
In [10]: #Bar Plot
sns.boxplot(x='status',y='mba_p',data=df)
plt.title('MBA Percentage VS Placement percentage')
plt.show()
```



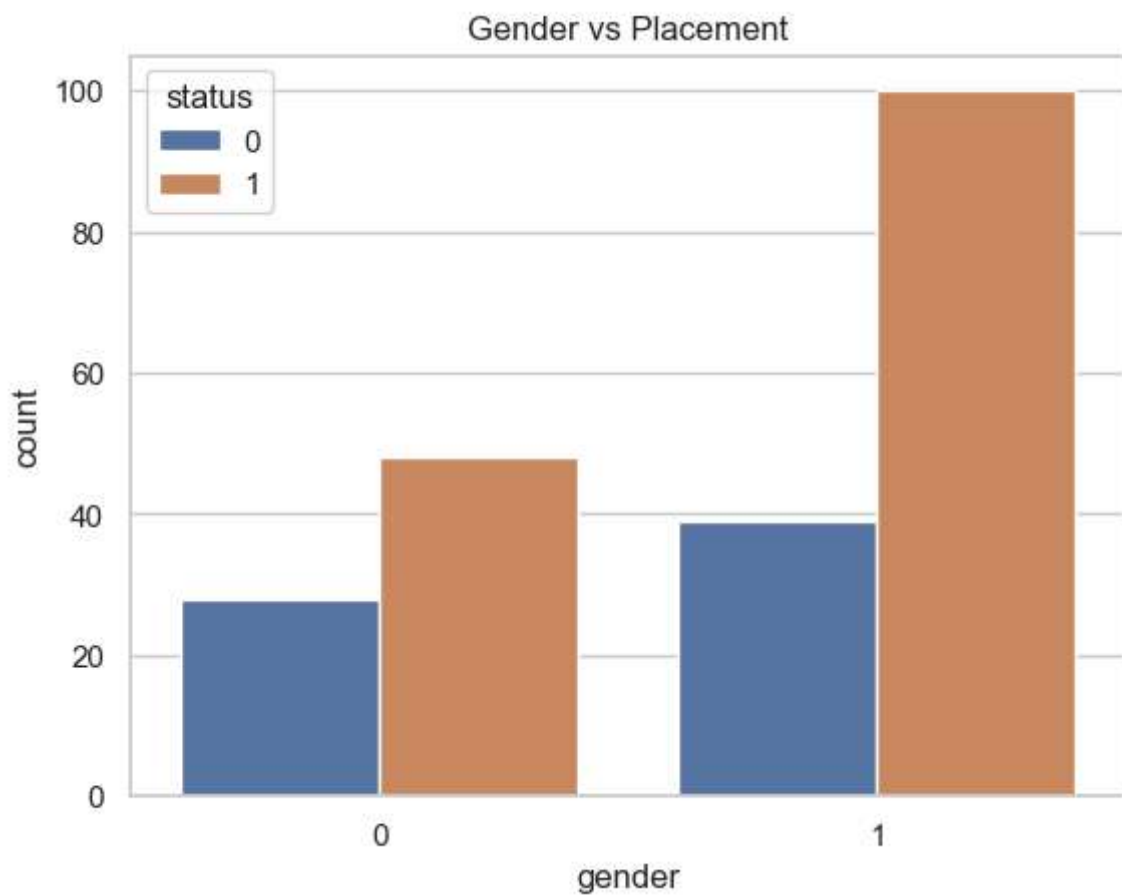
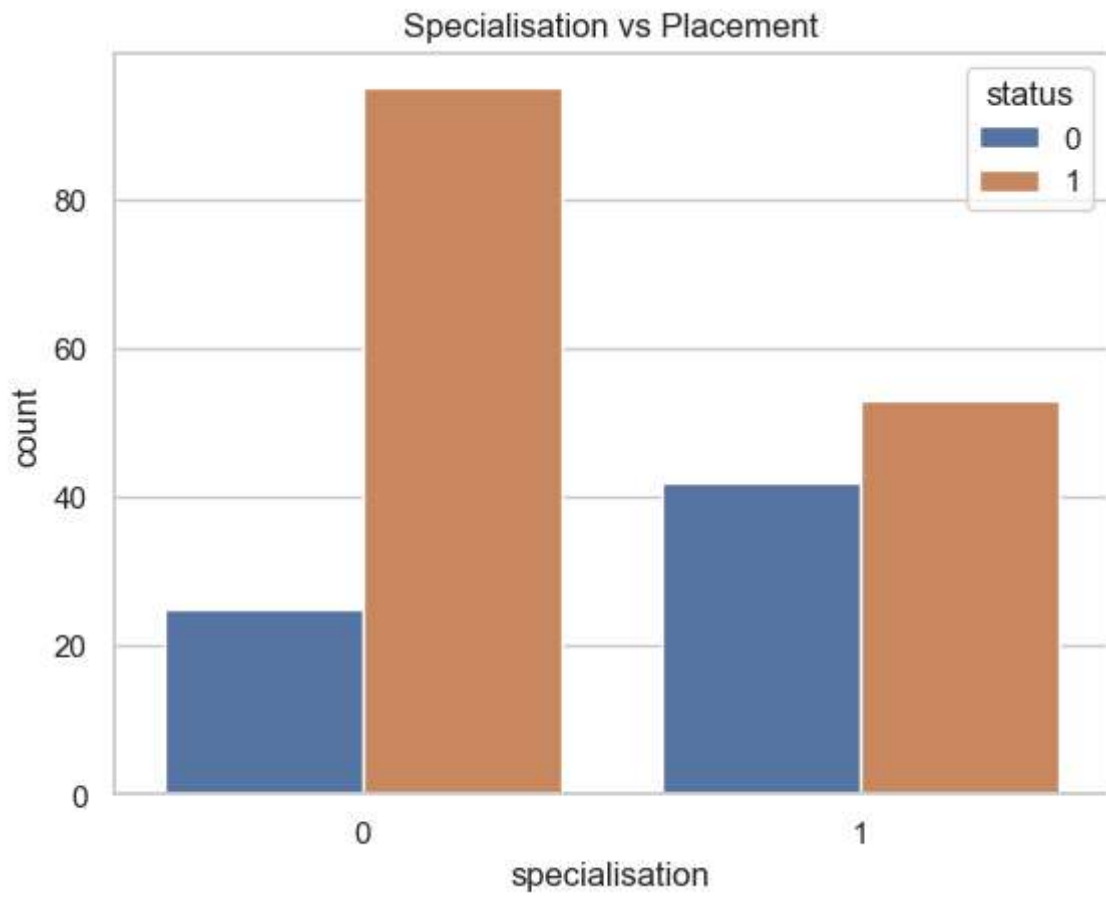
```
In [11]: #Correlation heatmap
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(),annot=True,cmap='coolwarm')
```

```
plt.title('Feature Correlation Heatmap')
plt.show()
```



```
In [12]: #Placement by specialisation
sns.countplot(x='specialisation',hue='status',data=df)
plt.title('Specialisation vs Placement')
plt.show()

sns.countplot(x='gender',hue='status',data=df)
plt.title('Gender vs Placement')
plt.show()
```



```
In [18]: df.to_csv("cleaned_data.csv", index=False)
```

```
In [5]: #Preprocess before modelling
import pandas as pd
df = pd.read_csv("cleaned_data.csv")
from sklearn.model_selection import train_test_split
X = df.drop("status", axis=1) # Features
y = df["status"]             # Target (0 = Not Placed, 1 = Placed)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta
```

```
In [6]: #Training multiple ML models

#Logistic Regression
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, accuracy_score
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
logreg = LogisticRegression(max_iter=500, solver='lbfgs')
logreg.fit(X_train_scaled, y_train)
y_pred = logreg.predict(X_test_scaled)
print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))

#Decision Tree
from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))

#Random forest
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
```

```
Logistic Regression Accuracy: 0.8837209302325582
      precision    recall  f1-score   support


     0       0.82      0.75      0.78         12
     1       0.91      0.94      0.92         31

 accuracy          0.88         43
  macro avg       0.86      0.84      0.85         43
 weighted avg     0.88      0.88      0.88         43
```


```
Decision Tree Accuracy: 0.8372093023255814
Random Forest Accuracy: 0.7906976744186046
```

```
In [7]: import matplotlib.pyplot as plt
```

```
accuracies = {  
    'Logistic Regression': 0.88,  
    'Decision Tree': 0.84,  
    'Random Forest': 0.79  
}  
  
# Bar chart  
plt.figure(figsize=(8,5))  
plt.bar(accuracies.keys(), accuracies.values(), color=['skyblue', 'lightgreen', 'salmon'])  
plt.ylabel('Accuracy')  
plt.title('Model Comparison')  
plt.ylim(0, 1)  
plt.show()
```




```
In [8]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay  
  
cm = confusion_matrix(y_test, y_pred)  
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=logreg.classes_)  
disp.plot(cmap=plt.cm.Blues)  
plt.title('Confusion Matrix - Logistic Regression')  
plt.show()
```



```
In [9]: from sklearn.metrics import roc_curve, auc

y_prob = logreg.predict_proba(X_test_scaled)[:,-1]
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'Logistic Regression (AUC = {roc_auc:.2f})')
plt.plot([0,1], [0,1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```

In []: